University of Glasgow | School of Computing Science

Exploring the possibility of accelerating Hurricane Simulator using Apache Spark framework

Andrej Hoos

School of Computing Science
Sir Alwyn Williams Building
University of Glasgow
G12 8QQ

Level 4 Project — March 24, 2016

Abstract

Large Eddy Simulator is a high resolution meteorological simulator designed to model flows over urban topologies. While the simulation has been accelerated with the use of OpenCL technology, there are still limits to the size of the simulation that can be run on a single computer.

This project tries to overcome this limitation by parallelising the simulation on a network cluster using Apache Spark framework. Suitability of multiple approaches to halo exchanges over network is explored and it is demonstrated that such exchange can be done using MapReduce frameworks and Apache Spark specifically. The simulation is rewritten in Java and uses Aparapi for OpenCL acceleration and Java Native Access to access Fortran routines in the original code. A number of simple test was conducted to determine the scalability and viability of this method. However, significantly decreased performance was observed when compared to the non-parallelised implementation due to framework overhead.

Education Use Consent

I hereby give my permission for this project to be shown to other University of Glasgow students and to be distributed in an electronic format. **Please note that you are under no obligation to sign this declaration, but doing so would help future students.**

Name: _____   Signature: _____

Acknowledgments

Contents

Chapter 1

Introduction

Simulating real world weather system is a difficult task, requiring a lot of computational power. One such simulator is The Large Eddy Simulator for the Study of Urban Boundary-layer Flows (LES) developed by Hiromasa Nakayama and Haruyasu Nagai at the Japan Atomic Energy Agency and Prof. Tetsuya Takemi at the Disaster Prevention Research Institute of Kyoto University [2] [3]. It simulates meteorological systems in an urban environment with building resolution. Originally LES is implemented in Fortran 77 with no explicit parallelization.

1.1 Motivation & Aim

As multi core computational devices are now commonplace, parallelisation has become a important consideration when implementing simulations. Therefore LES has been ported to Fortran 95/OpenCL implementation by Wim Vanderbauwhed [7]. Using this implementation, LES can be run parallelised on any OpenCL compatible computational device, which shows significant improvements in speed. While these improvements can reduce the time needed to run single simulation run or increase the size of the spatial domain of the simulation, there are limits to both. One of these limits is number of computational units the device has, limiting the extent of parallelization. The other is memory available on the device.

Using a network cluster of computers can, in theory, help overcome these limits to some extent. However, splitting work onto multiple nodes in a cluster introduces a issue of communicating important information over network. Specifically, in case of LES, data that is at the edge of the spatial domain of each node needs to be exchanged in so called halo exchanges.

Recently, there has been a rise in a new computational frameworks focused on working in distributed environment. One of these frameworks is Apache Spark which is based on the MapReduce parallel execution paradigm.

The aim of this report is to explore the possibility of using the Apache Spark in Java to parallelise LES, or any algorithm relying on halo exchanges in network distributed environment.

1.2 Outline

Further background will be described in chapter 2, including information on LES, Apache Spark and OpenCL implementation used. Implementation of halo exchanges in Apache Spark will be discussed in chapter 3 and

chapter 4 will explain the process of porting LES code from Fortran to Java. Evaluation of the achieved performance can be found in chapter 5 and chapter 6 contains final remarks and notes about future work.

Chapter 2

Background

2.1 Large Eddy Simulator

The Large Eddy Simulator for the Study of Urban Boundary-layer Flows developed by Disaster Prevention Research Institute of Kyoto University is a high resolution meteorological simulator designed to model turbulent flows over urban topologies. It achieves this by using mesoscale meteorological simulations and building resolving large-eddy simulations. It also uses The Weather Research and Forecasting Model [1] to compute the wind profile as an input for the simulation. The simulation is split into multiple subroutines, however, the most important one is the Poisson Equation solver. This is the most computationally intensive subroutine and it uses successive over relaxation to solve a linear system of equations. Each time step of the simulation needs to execute following subroutines in a sequence:

velnw	Update velocity for current time step
bondv1	Calculate boundary conditions (initial wind profile, inflow, outflow)
velfg	Calculate the body force
feedbf	Calculation of building effects
les	Calculation of viscosity terms (Smagorinsky model)
adam	Adams-Bashforth time integration
press	Solving of Poisson equation (SOR)

The simulator was originally written in Fortran 77. The current version was ported to Fortran 95 and accelerated using OpenCL in work by Wim Vanderbauwhed [7].

LES runs over spatial domain represent by a 3D grid. Size of 150x150x90 is used for the most of this report, but the size can be varied in x and y direction. By increasing the size of the grid in x and y directions, LES can be run with higher resolution, or it can used to cover larger area. However doing this increases the time needed to complete the simulation and more importantly there is a limit to how big the grid can be so that it fits inside the RAM.

2.1.1 Parallelisation

Using a cluster of computers, we can split the domain such that every node in the cluster runs LES on a small portion of the domain. This allows to run LES on much larger domain than what would be possible on a single
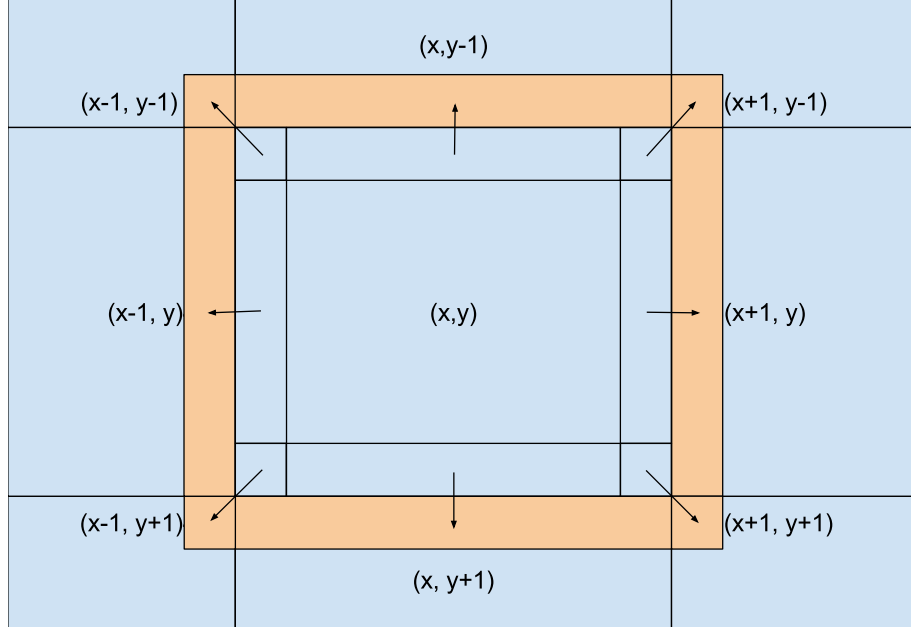
---

[1] http://www.wrf-model.org

Figure 1: Node (x,y) needs access to the edges of all its neighbours

computer. To run LES correctly in this scenario, each node needs access to the edge data belonging to each of its eight neighbours (figure 1).

Message Passing Interface (MPI) [2] is considered to be most commonly used for communication in distributed memory environments. Because of this, LES has already been parallelised using MPI by Gordon Reid as his MSci project at University of Glasgow [5] and others [4]. Although MPI is considered standard and achieves good performance, it isn't easy task to parallelise algorithms using MPI, especially for scientists with minimal experience with Computing Science. Therefore it is important to investigate the possibilty of using frameworks with higher level abstraction and that are easy to deploy onto a standard network cluster.

2.2 MapReduce parallel processing

MapReduce is a programming model used mainly for processing large quantities of data and is intended for network clusters [1]. Program using this model consists of Map and Reduces methods which are automatically parallelised. MapReduce was originally name of a technology by Google, but is now used for the model in general. MapReduce libraries and frameworks are now available in many programming languages.

MapReduce model usually operates on key-value pairs and can be described by three main steps:

- **Map**: **map()** function is applied to all of the pairs, creating a collection of new key-value pairs

- **Shuffle**: the output data is redistributed across network such that pairs with the same key are located on the same node

- **Reduce**: the output data is processed by applying the **reduce()** function to the output pairs

---

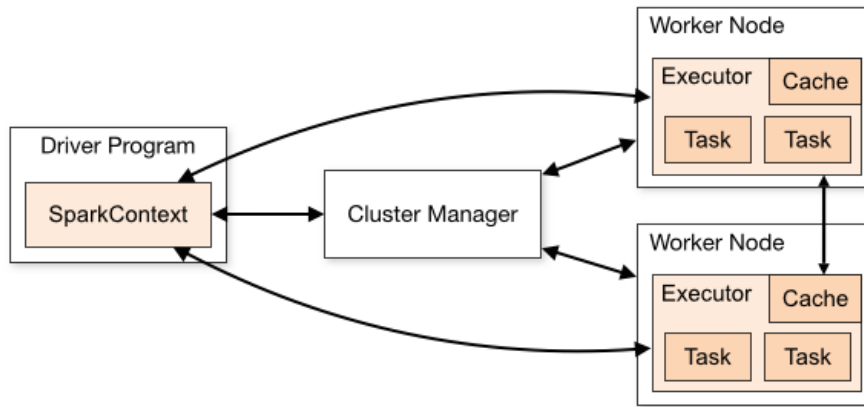[2]http://www.mcs.anl.gov/research/projects/mpi/mpi-standard/mpi-report-2.0/mpi2-report.htm

Figure 2: Apache Spark cluster components overview [4]

Advantage of using MapReduce is that only map() and reduce() functions need to be written and the shuffle step is handled automatically by the library or framework and is usually optimised. In addition, most MapReduce frameworks provided more data manipulation methods such as grouping or joining.

2.2.1 Apache Spark

Apache Spark[3] is one of these frameworks and is available Java, Scala, Python and R. It claims to be faster over alternatives such as Apache Hadoop, because it allows MapReduce to run in memory as opposed to reading and writing inputs and outputs solely from disk. This, with combination that is the increasing community and high-level abstraction makes this framework suitable for the purposes of this project.

Every Spark program consists of a driver program and a worker node(s). When Spark is used in cluster mode, a cluster manager is also required (figure 2). Cluster manager allocates resources across nodes and is usually referred to as Spark master in this report. The driver program Once the driver program connects to the cluster, it acquires executors on the worker nodes and starts sending tasks for them to execute.

The driver program runs the user's main function and allows to run parallel operations on the cluster. The main abstraction provided by spark is represented by resilient distributed dataset (RDD). RDD is a collection of elements which is automatically distributed across the cluster that can be operated on in parallel.

Listing 1 shows an example Spark program which loads a text file and creates a RDD containing words from the file. It then uses `mapToPair` to create key value pairs, where key is a word a value is count and then produces counts of each word by using `reduceByKey`.

2.3 OpenCL

OpenCL[5] is a framework that allows to writing parallelised software on various computational platforms, such as CPUs, GPUs and hardware accelerators. OpenCL uses own programming language based on C99 for executing programs called kernels on these platforms. These kernels can be run across multiple devices with almost no

---

[3] http://spark.apache.org/
[4] http://spark.apache.org/docs/1.4.1/cluster-overview.html
[5] https://www.khronos.org/opencl/

```
1  SparkConf conf = new SparkConf().setAppName(appName).setMaster(master);
2  JavaSparkContext sc = new JavaSparkContext(conf);
3
4  JavaRDD<String> textFile = sc.textFile("...");
5  JavaRDD<String> words = textFile.flatMap(new FlatMapFunction<String, String>() {
6    public Iterable<String> call(String s) { return Arrays.asList(s.split(" ")); }
7  });
8  JavaPairRDD<String, Integer> pairs = words.mapToPair(new PairFunction<String, String, Integer>() {
9    public Tuple2<String, Integer> call(String s) { return new Tuple2<String, Integer>(s, 1); }
10 });
11 JavaPairRDD<String, Integer> counts = pairs.reduceByKey(new Function2<Integer, Integer, Integer>() {
12   public Integer call(Integer a, Integer b) { return a + b; }
13 });
14 counts.saveAsTextFile("...");
```
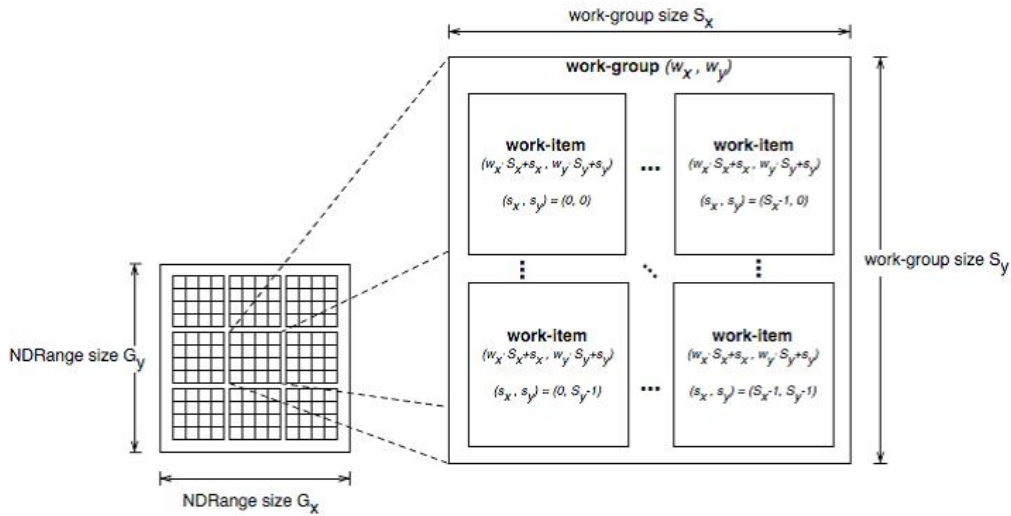
Listing 1: Example Spark program

Figure 3: Work structure in OpenCL [6]

modification. The differences in code for different devices are usually because of different feature set (e.g. double floating point precision support), or different performance of specific operations.

2.3.1 Structure overview

OpenCL architecture consists of a host and a computational device connected to the host. The host is usually the CPU and it handles the initialisation, allocating memory, queuing work and harvesting of results from the computational device. The computation devices usually have one or more compute units and each compute unit consists of multiple processing elements. A single kernel execution usually runs on multiple processing elements in parallel.

A work item is the smallest control flow unit in the OpenCL architecture. A work group is a collection of related work items, which execute the same kernel and share local memory. The size of the work group is limited by the underlying device and can usually consist of up to three dimensions.

When executing the kernel, the host code specifies the global work size. The host can specify the size of the work groups the work is divided into, or it can be left to implementation do decide automatically. The kernel is

---

[6] https://software.intel.com/sites/landingpage/opencl/optimization-guide/Basic_Concepts.htm

then executed on each work item, which can determine its portion of work by accessing its global or local id. The figure ... shows the overall work structure in OpenCL.

2.3.2 Kernels

Kernels are OpenCL programs executed in parallel on a computation device. The language that is used to write kernels is based on C99, but adapted to fit the OpenCL device model. It replaces the C standard version library with custom set functions, which contain additional features, such as math functions. When using pointers, programmer can use region qualifiers annotations, such as `__global` or `__local`, which specify a level in memory hierarchy. Besides standard types like `float` or `int`, OpenCL defines vector type with fixed length with size of 2, 4, 8 or 16. Operations on these vector types are intended to map onto vector instructions on the device. The entry point functions are marked with `__kernel` which signals that it can be called by the host code.

The host code is responsible for selecting the device, loading the kernel code and building the program it using the selected device. The program can be than queued, which also involves specifying the global work size, and optionally the local work size.

2.3.3 Aparapi

Aparapi[7] is an API that allows programmers to write OpenCL capable parallel programs in Java. Aparapi was originally an AMD product [8], but was released to open source in 2011. It is capable of converting Java bytecode into OpenCL programs that can be executed on computational devices. Aparapi also provides abstraction for selecting devices and queuing programs.

Listing 2 shows an example how Aparapi works. The example creates two vectors, `a` and `b` and sums them into vector `sum`. `a` and `b` are represented by two arrays, which are initialised with random values. The OpenCL program is created by using a `Kernel` class with overloaded `run` method and can be enqueued and run using `execute` method. The underlying OpenCL code that is used to create the OpenCL program can be seen in listing 3.

2.4 Related work

Parallelised version of LES has been created by Gordon Reid [5] using MPI and Glasgow Model Coupling Framework (GMFC). The work creates MPI implementation for both shared- and distributed-memory systems. GMFC implementation is created for shared-memory system. The implementation is evaluated using constant global area as well as using expanding area test with each node having area of the same size. Both MPI and GMCF have been shown to scale well using both evaluation methods.

Another work to Parellise LES is work by Raasch and Schrter [4]. The work uses different LES implementation from the one used in this project. Parallelisation using MPI is shown to speed up the simulation on both small domain and large domains.

Using Apache Spark for halo exchanges seems to be not well researched use case, although one work [8] has used Spark to solve a stencil problem computation. The work implements Jacobi Stencil Codes algorithm which requires exchange of boundary data. Three different methods are evaluated. Method using broadcast variables,

---

[7]https://aparapi.github.io/
[8]http://developer.amd.com/tools-and-sdks/opencl-zone/aparapi/

```
1  import com.amd.aparapi.Kernel;
2  import com.amd.aparapi.Range;
3
4  public class VectorAdd
5  {
6      public static void main( String[] args )
7      {
8          final int size = 50000000;
9
10         final float[] a = new float[size];
11         final float[] b = new float[size];
12
13         for (int i = 0; i < size; i++) {
14             a[i] = (float) (Math.random() * 100);
15             b[i] = (float) (Math.random() * 100);
16         }
17
18         final float[] sum = new float[size];
19
20         Kernel kernel = new Kernel(){
21             @Override public void run() {
22                 int gid = getGlobalId();
23                 sum[gid] = a[gid] + b[gid];
24             }
25         };
26
27         kernel.execute(Range.create(size));
28     }
29 }
```

Listing 2: Example Aparapi program that sums two vectors

```
1  typedef struct This_s{
2      __global float *val$sum;
3      __global float *val$a;
4      __global float *val$b;
5      int passid;
6  }This;
7  int get_pass_id(This *this){
8      return this->passid;
9  }
10 __kernel void run(
11     __global float *val$sum,
12     __global float *val$a,
13     __global float *val$b,
14     int passid
15 ){
16     This thisStruct;
17     This* this=&thisStruct;
18     this->val$sum = val$sum;
19     this->val$a = val$a;
20     this->val$b = val$b;
21     this->passid = passid;
22     {
23         int gid = get_global_id(0);
24         this->val$sum[gid]  = this->val$a[gid] + this->val$b[gid];
25         return;
26     }
27 }
```

Listing 3: OpenCL code generated by the Aparapi example

method that shares the boundary data using CassandraDB and method that uses boundary RDDs. Method with broadcast variables is shown to scale well up to around 100 nodes, while the other methods can scale up to 500 nodes, with the boundary RDDs approach providing better performance. Unfortunately the work does not describe in detail the implementation details of this approach.

Chapter 3

Halo exchanges in Spark

The main challenge of this project is to construct viable halo exchange algorithm, using data manipulation methods provided by Apache Spark. Multiple possible algorithms were considered and their implementation details, advantages and drawbacks are discussed in this chapter.

3.1 Game of Life as a proof of concept

To test the considered algorithms, a implementation of Conway's Game of Life was used. This was because of its simplicity and becasue it resembles the simulator in that in needs data from neighbours to be transfered in order to be parallelised in distributed environment.

The Game of Life is a cellular automaton and is run on a 2D grid of cells. To compute the transition of a single cell, state of each of its eight neighbours are needed. The transitions are defined as:

- Any live cell with fewer than two live neighbours dies.

- Any live cell with two or three live neighbours lives on to the next generation.

- Any live cell with more than three live neighbours dies.

- Any dead cell with exactly three live neighbours becomes a live cell.

We can group the cells into same size rectangle shaped chunks and compute the transitions of each chunk on a separate node of the cluster. However, now the nearest edges of neighbouring chunks need to be exchanged before each transition, in order for the edge cells to know the state of all of their neighbours. This is the same case in the LES and is therefore suitable as a proof of concept.

3.2 Broadcast variables

The simplest way exchange halos is using broadcast variables. Spark allows to create a broadcast variable in the driver program, which is the shipped to worker nodes if needed. However the limitation is that only the driver program can create broadcast variables (i.e. they cannot be created in the worker code) and therefore the workers

Figure 4: Diagram of the the algorithm using broadcast variables

would need to send their edge data to the driver program, which would contruct the appropriate halos and create a broadcast variable(figure 4).

This means that the halos would be transferred to their destination always via the node which hosts the driver program, and therefore the algorithm would do twice as much network transfers as the ideal case where the halos are transferred to their destination directly. Therefore this approach was deemed unsuitable.

3.3   MapReduceByKey

Another way is to use a combination of map and reduce steps (figure 5). To represent the chunks, we use Spark Pair Collection of pairs of key and chunk (`JavaPairRDD<String, Chunk>`). The key represents the coordinates of the chunk.

In the first stage we transform this collection into another collection which contains the chunks from the original collection together with all neighbours using map (`flatMapToPair`). The keys of the neighbours key value pairs are the keys of their respective target chunk.

The second step then does reduce on this new collection using `reduceByKey`. The reducer function takes a chunk and a neighbour and reduces them into a chunk with the neighbour added to it, which results in a collection of chunks which are larger by one cell on each side.

The last step is a map step, which computes next transition on each chunk, returning collection of transformed chunk which are the original size.

There are, however, some drawbacks to this approach. The first is, that we cannot control the exact order in which these operations are executed. For example, Spark might decide to execute the first map step on first
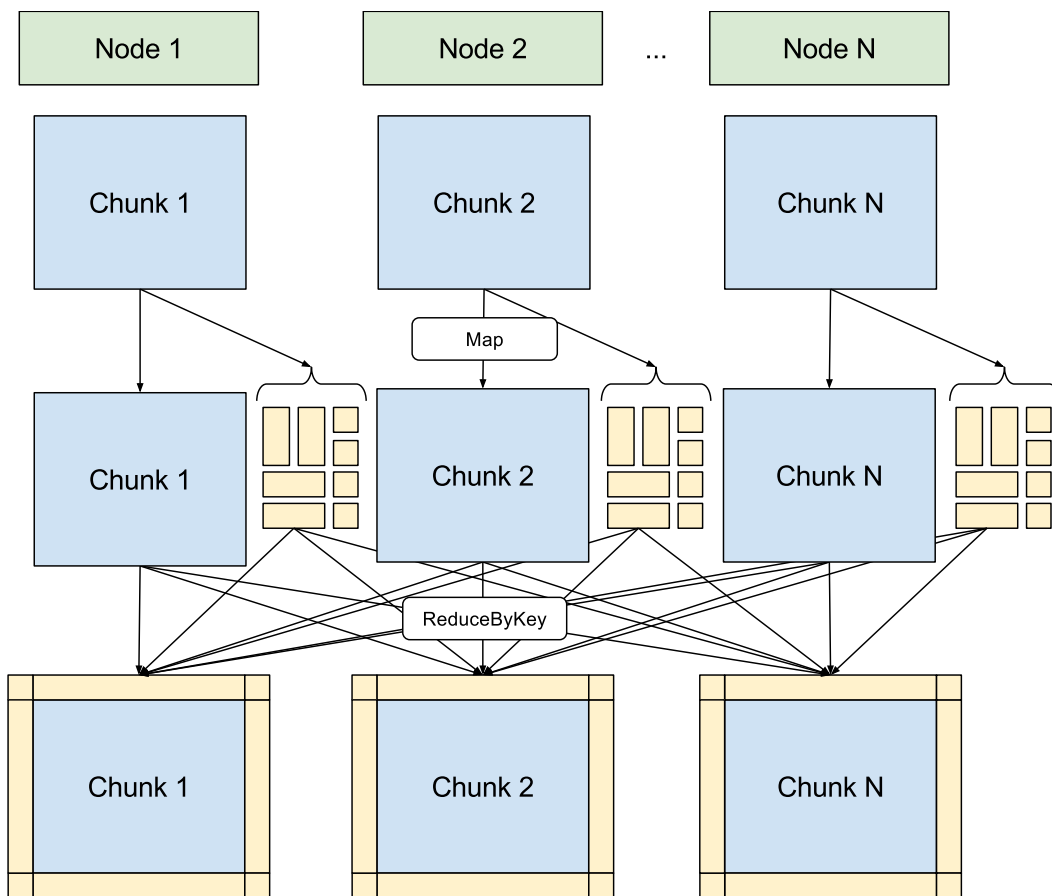
11

Figure 5: Diagram of the the MapReduceByKey algorithm

two chunks (assuming there are more than two) and run the reduce step on the results. Then, the first map step might be run on the rest of the chunks. The second issue is that we cannot control, on which node is each step executed. This two issues result in all of the chunks and not only neighbours are transferred over network, which is unacceptable for the purposes of this project.

3.4 MapGroupByKey

To resolve one of the issues mentioned above, it is possible to replace the reduce step with group step (figure 6). This uses `groupByKey` to group chunk together with neighbours that share the same key. The grouping function takes the chunk with its neighbours and combines them into a single chunk, such that the result is the same as after the reduce step in previous section.

Using this method the issue of the order of execution is solved, since the group can be executed only after all of first map stage has been executed. The second issue of not being able to control where is each stage executed still remains and therefore the issue of transferring all chunk over network remains as well.

3.5 MapCogroup

Method described in this section was inspired by chapter 4 in *Learning Spark* [9]. It uses custom partitioner and two separate collections for chunks and neighbours to prevent Spark from moving chunks around the network. It uses combination of map and cogroup steps.

The first map operates on the same key value pair collection as above. However, now it transforms the collection into separate key value pair collection of neighbours, instead of creating one mixed one as was done previously (figure 7).

In second step, the `cogroup` transformation is very similar to the grouping transformation from the MapGroupByKey method, with the difference that in this case it operates on two collections instead of one. The result is the same as in previous sections.

In addition, the collection containing chunks is partitioned using custom partitioner. The partitioner takes the key, which represents coordinates and outputs index of node on which the chunk should be physically present. Using this partitioner spark shuffles only the other collection, which contains neighbour data, over network and keeps the chunks collection in place.

3.6 Caching and checkpointing

Spark uses lazy execution when running programs stages. This means that transformation on a collection are saved in its lineage are executed only after a call that materialises the collection on the driver program is called on it. Examples of such calls are `count` that returns the number of items in the collection, `collect` which returns array of all items in the collection or `reduce`, which, as opposed to `reduceByKey` reduces the collection into single item.

This introduces some issues. First issue is that often, when resolving stages that do network shuffle, Spark executes the parent stages more than once. This is considered standard Spark behaviour and is solved by caching
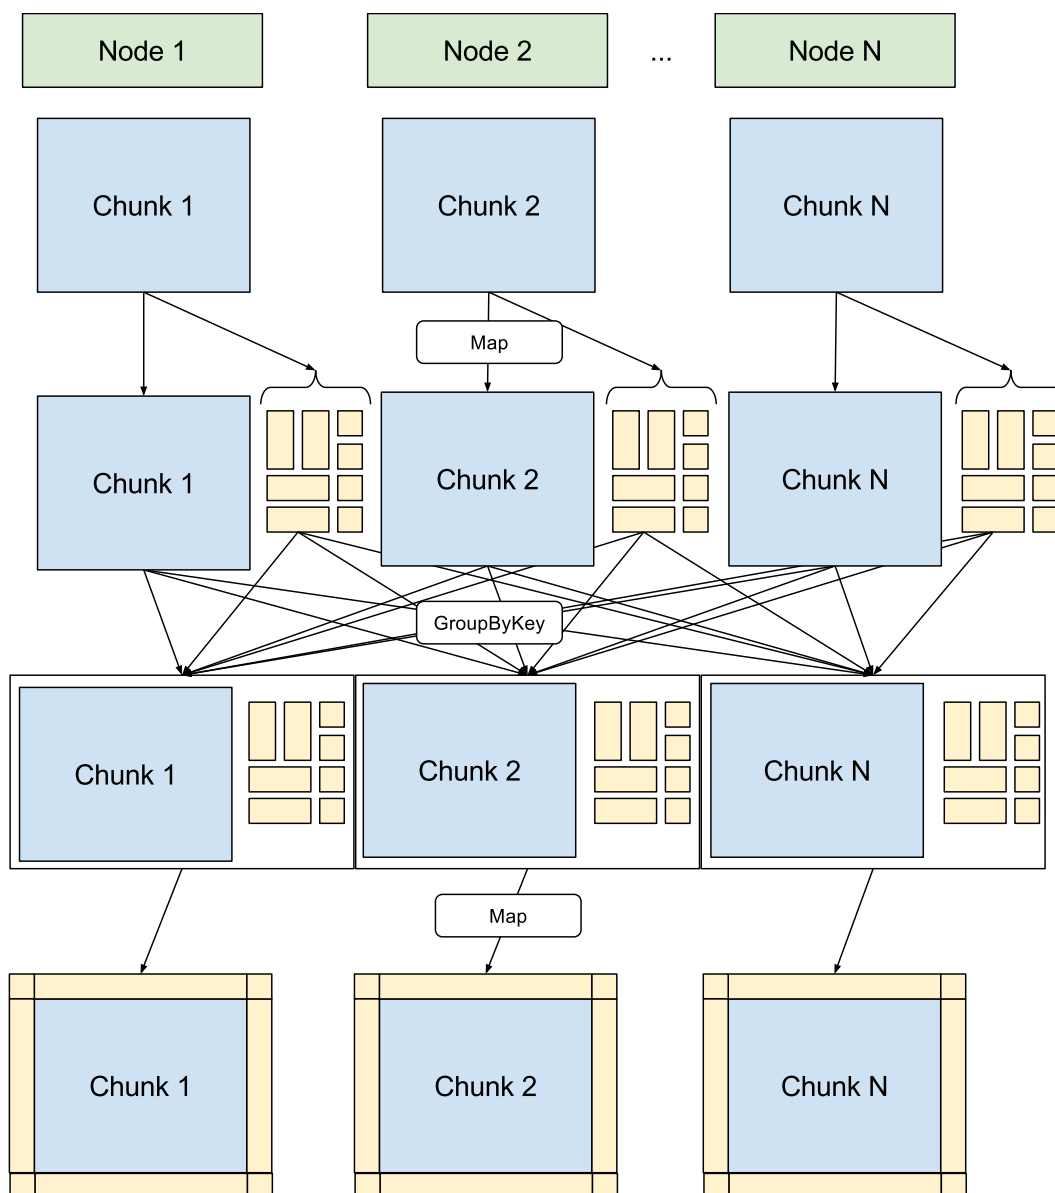
Figure 6: Diagram of the the MapGroupByKey algorithm
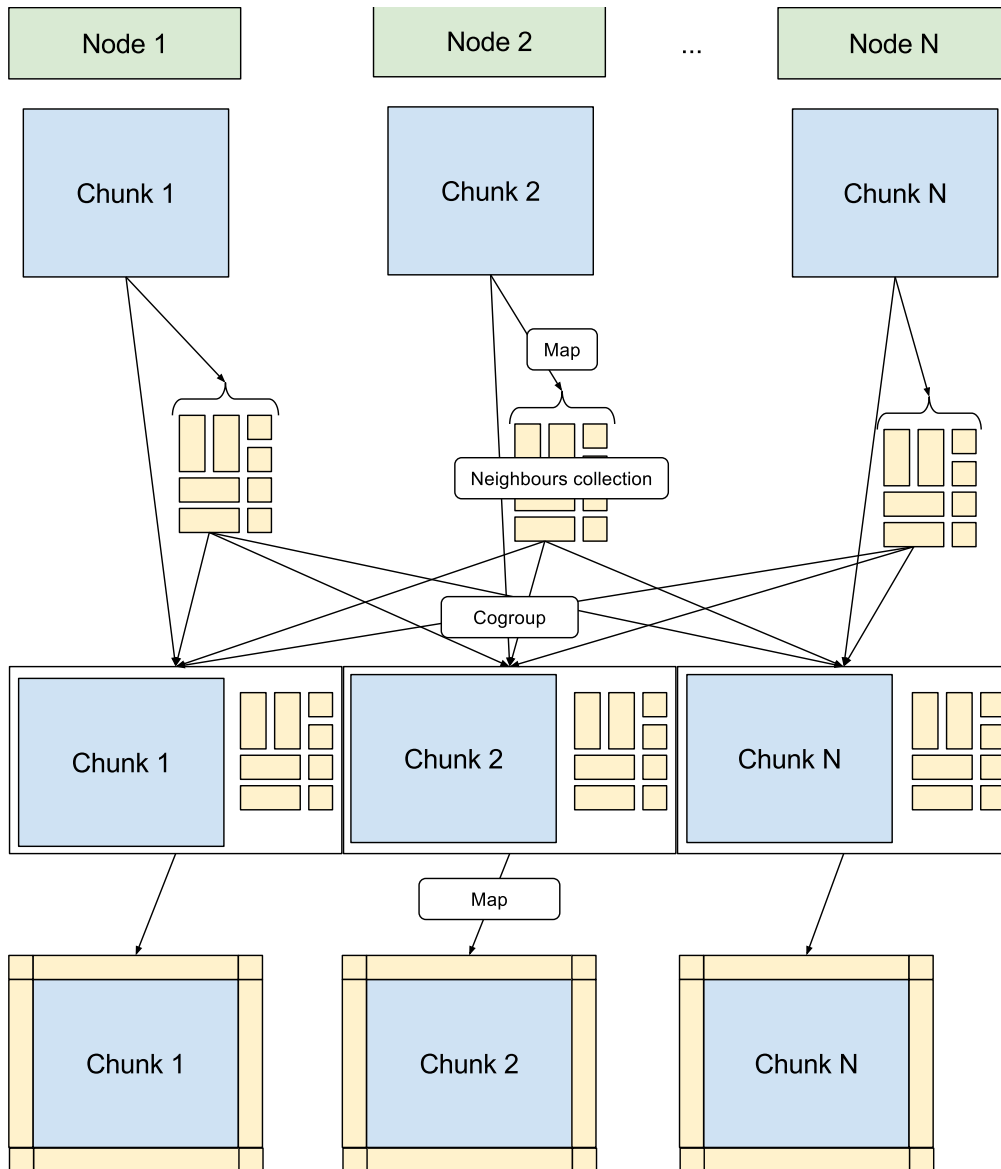
Figure 7: Diagram of the the MapCogroup algorithm

the parent collection using `cache` method, which also marks the collection to be explicitly reused when executing, child stages. However, it was found that Spark keeps more cached collections than necessary, and therefore workaround was devised. Using this workaround, first, new collection is created, then cached and materialised using `count` and then `unpersist` is called on the old collection to remove it from the cache.

Another issue is lineage length. In iterative algorithms, such as the ones described in this report, the lineage that Spark keeps in order to be able to restore execution in the case of failure can get quite long. This results in significant slowdown of Spark task resolution and eventually leads to stack overflow errors. One possible solution is to increase JVM maximum stack memory, but the standard practice is to `chekpoint` collections. This serialises the collection to disk and forgets its lineage.

Chapter 4

LES in Java

The second challenge this chapter deals with is reimplementing LES in Java. To do this, the OpenCL host code written in Fortran needs to be rewritten Java. This code consists mainly of a `switch` construct, which executes different routines in the OpenCL kernel. Next methods for initialising variables, constructing and deconstructing halos need to be created. Finally, the newly written Java OpenCL host code needs to be integrated with Spark.

4.1 JNA with Fortran

LES works with about 40 different variables that store states of different systems, wind velocity, pressure, halos and grid constants. These variables need to be initialised before the execution of the simulation. This process is covered in a substantial part of the Fortran source code.

Since accessing the existing Fortran subroutines is a viable strategy [1], it was decided to write a simple Fortran subroutine, which groups all subroutines that initialise the needed variables. The Fortran code is compiled into a shared library, and can be accessed from Java using Java Native Access [2].

4.2 Aparapi for Unconventional Cores

To write the OpenCL host code Aparapi for Unconventional Cores (Aparapi-ucores) [6] [3] is used as an underlying library that facilitates communication with the OpenCL kernel. It is a fork of Aparapi library that allows running OpenCL code on unconventional cores such as FPGAs. In addition it allows running OpenCL kernels from precompiled binaries.

Aparapi allows to write Java code which is then translated into OpenCL kernel code and executed. However, since the LES kernel code is already written, this is counterproductive. It is impossible to disable this behaviour and therefore a workaround is necessary. Although, with binary flow, which is provided by Aparapi-ucores it is possible to execute compiled binary of the provided kernel instead of kernel generated by Aparapi, Aparapi will expect the same method signature as the one the generated kernel has. Therefore it is needed to provide Aparapi with dummy Java code, that translates into correct OpenCL method signature. This can be achieved by using all parameters in correct order, the parameters need to have correct datatype and parameters that are supposed to be

---

[1] `http://www.javaforge.com/wiki/66061`

[2] `https://github.com/java-native-access/jna`

[3] `https://gitlab.com/mora/aparapi-ucores`

17

written to by kernel need to be marked as such by assigning a value into them. Example dummy code is in the listing below:

```
1  float[] p2;
2  float[] uvw;
3  float[] uvwsum;
4  float[] fgh;
5
6  @Override
7  public void run() {
8    float sum = p2[0] + uvw[0] + uvwsum[0] + fgh[0];
9    uvw[0] = sum;
10   uvwsum[0] = sum;
11 }
```

The corresponding kernel method signature as generated by Aparapi:

```
1  __kernel void run(
2    __global float *p2,
3    __global float *uvw,
4    __global float *uvwsum,
5    __global float *fgh,
6    int passid
7  )
```

In the example above, four kernel parameters are used. All are arrays of floats. Note that the order in which they are listed in the sum in the first line of the `run` method, is the same as the order of the parameters in the signature. In addition the `uvw` and `uvwsum` variables were assigned into, which marks them as kernel-writable inside Aparapi. Also, important to note, is final parameter `passid`, which is being automatically inserted by Aparapi, and contains the id of the current kernel execution.

Due to the nature, there are some modification that need to be done to the OpenCL kernel, before compiling into a binary. First, the `passid` needs to be added to the parameter list. Second, Aparapi does not support passing non-constant scalar fields into the kernel as parameter. As a workaround to this, scalar fields can be represented as simple single item 1D arrays. Finally the name of the kernel method is restricted to the name `run`. As of writing of this report, however, this was found not to be true and it is possible to use different names, but the written code does not reflect this.

By default, Aparapi copies all parameters to and from kernel before and after every kernel execution. This is due to large number of parameters in LES undesirable. Therefore, for iterative applications, it is recommend to enable explicit Aparapi mode, which then requires the programmer to call `put` and `get` methods to explicitly write and read parameters to and from the kernel.

4.2.1 Device specific OpenCL binaries

Since OpenCL is supposed to be compatible with wide range of devices, the same OpenCL code can be executed on any OpenCL device. This is not true, however, for the compiled binaries, which are device specific. This means that new binaries need to be created whenever the simulation is to be run on a new device. There is no tool available that allows easily compile kernel into binaries, but OpenCL specification describes a method that allows to retrieve the compiled binary code [4]. Therefore it was decide to create a simple compilation tool [5] which allows the user to specify a `.cl` kernel file and device with which to compile the kernel file.
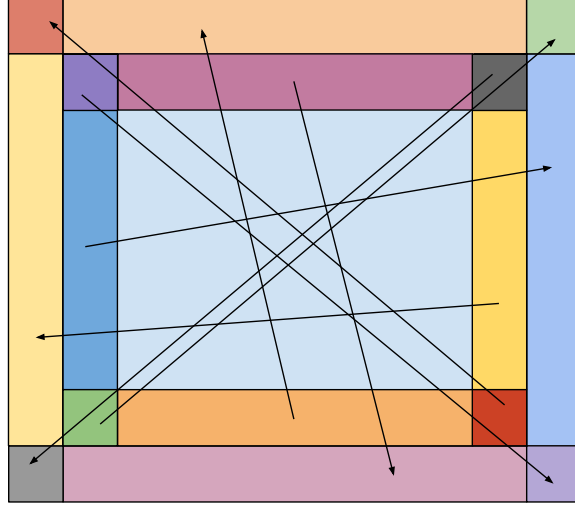
---

[4] https://www.khronos.org/registry/cl/sdk/1.0/docs/man/xhtml/clGetProgramInfo.html
[5] https://github.com/adikus/cl-compile

Figure 8: Diagram of halo change on a single node

4.2.2 Memory alignment

Aparapi-ucores contains functionality which copies the contents of the parameters into a memory aligned location. This is done because some devices such as FPGAs do not interact well with unaligned memory locations, and locations in JVM heap might not be aligned. This functionality is currently experimental, and allocates new memory for every kernel-writable parameter every kernel execution. In addition it copies the content of parameters from the original JVM memory location to this new location and back every kernel execution for every parameter. This results in a significant slowdown in execution, which is undesirable.

This functionality can be turned of, since in theory, using unaligned parameters should not pose any issues when running the kernel on CPUs and GPUs. However, after switching this functionality it was found, that the kernel no longer runs properly on Intel CPU. The reason remains, as of writing of this report, unknown, and therefore it was decided to modify the underlying Aparapi-ucores C++ source code to create a workaround that deals with the excessive memory copies. The workaround allocates the aligned memory location only once and copies the memory from or to it only when the `get` or `put` Aparapi methods are called.

4.3 Halo construction and deconstruction

Since LES has already been parallelised using MPI [5], functionality for reading and writing halos is already implemented inside the OpenCL kernel code. This format is a single 1D for each halo and therefore needs to be deconstructed before sending and constructed again after receiving.

Since the spatial domain in this project is considered to be wrapped such that bottom connects to top and left connects to right, this halo construction and deconstruction can be demonstrated on a single node (figure 8).

4.4   LES in Apache Spark

As a final part of the project implementation, working LES Aparapi implementation needs to be integrated into Spark. MapCogroup described in chapter 3 is used as the halo exchange algorithm.

Each time step of the simulation is divided into 11 steps, with each step executing OpenCL kernel once. The 8th step (`PRESS_SOR`) is a exception. In this step the successive over relaxation is run, which requres multiple iterations of kernel executions. Normally the over relaxation is run until the simulation state converges, but in this project it was decided to use constant number of 50 iterations.

Each kernel execution needs to be accompanied be a halo write beforehand and halo read afterwards, both of which are an extra kernel execution. After each group of halo write, simulation step and halo read the appropriate halos should be exchanged over network. In addition some stages do a reduction over the whole domain, which means that a reduction over the nodes needs to be done using Spark.

Table 1 contains the summary of each stage.

#	Stage	Halo exchange, reduction
1	`VELNW__BONDV1_INIT_UVW`	`p`, `uvw` and `uvwum` halo exchange
2	`BONDV1_CALC_UOUT`	`uvw` halo exchange
3	`BONDV1_CALC_UVW`	`uvw` halo exchange
4	`VELFG__FEEDBF__LES_CALC_SM`	`uvw`, `uvwsum`, `fgh`, `diu` and `sm` halo exchange
5	`LES_BOUND_SM`	`sm` halo exchange
6	`LES_CALC_VISC__ADAM`	`fgh` and `fgh_old` halo exchange
7	`PRESS_RHSAV`	`rhs` and `fgh` halo exchange + reduction
8	`PRESS_SOR`	`p` halo exchange, there are three kernel executions for each SOR iteration, which means three halo exchanges per iteration + one reduction per iteration to determine the convergence value
9	`PRESS_PAV`	`p` halo exchange + reduction
10	`PRESS_ADJ`	`p` halo exchange
11	`PRESS_BOUNDP`	`p` halo exchange

Table 1: Summary of halo exchanges and reductions for each LES stage in one time step

Spark collection that contains the kernel hosts consists of key value pairs. The key is represented by an integer which has value between 0 and X*Y-1, where X and Y are the size of the node grid. For example, if X is 2 and Y is 3, it means we intend to run the simulation on cluster with 6 worker nodes, and each node is running a simulation on a 150x150x90 size domain. The value in the key value pair is represented by a subclass of Aparapi Kernel class.

The neighbours are represent by a Neighbour class which encapsulates the underlying data array and the direction from which it came (N,S,W,E,NW,NE,SW,SE).

Finally, helper methods have been created for executing stage of simulation, exchanging halos and reductions. This means that a execution of a single simulation stage can be written similarly to:

```
// 7th stage - PRESS_RHSAV
kernels = executeKernelStep(kernels, States.PRESS_RHSAV);
kernels = HaloExchanger.exchangeHalos(kernels, "rhs,fgh", ip, jp, kp, X, Y);
kernels = divisionReduction(kernels);
```

4.4.1 Lineage & Checkpointing

The fact that single timestep of the simulation needs to be represented by large number of jobs (10 stages * 2 [execution + halo exchange] + 50 SOR iterations * 3 * 2 = 320; this is just an approximation) means that lineage grows very quickly. With the recommended rate of chekpointing, which is once per 100 jobs, three checkpoints are created every time step of the simulation. Since one time step should not take longer than couple of seconds, this means that large quantity of data is serialised to disk over the run of program.

This can create problems when trying to execute the simulation in environment with restricted disk space. Therefore another workaround was devised, which all previous checkpoints before creating a new one. Checkpoints in Spark are used to recover from a crash of a node or a similar scenario. However in our case, recovering is not as simple as deserialising a class from disk. Further work would need to be made to initialise the OpenCL kernel and load the recovered state of parameters into it, and was deemed to be out of the scope of this project. Therefore the checkpoints are completely useless and can be deleted without worry.

Chapter 5

Evaluation

Evaluation in this chapter was done to determine, whether the approach to parallelisation developed in this project is a competitive alternative to MPI.

Small number of tests was run without Spark to determine the true runtime of the simulation without overheard caused by Spark.

The final tests were done on a network cluster, with number of nodes varying from one to eight. However only very small number of tests was conducted, due to project timing issues (too close to deadline). Under normal circumstances, tests with higher number of nodes and different number of time steps would be run. In addition, all tests would be run multiple times.

5.1 Architecture

All tests were done on GPG cluster of CS department at University of Glasgow [1]. The cluster consists of 20 nodes, each having 64GB of RAM and 2 Intel Xeon processors [2]. Three nodes of the cluster are available to public and therefore only the rest 17 are suitable for running experiments.

5.2 Baseline tests

Two baseline tests were done. First runs LES as it would be run without any parallelisation over network. Second does the same, but adds halo writes and reads as well as simple halo exchanges as show in figure 8. Another test was done using Spark local mode. In this mode both the driver program and Spark worker work in a single container. Each test was run twice and all tests were run using 100 time steps. The results measured using the `time` command can be seen in table 2. Commands used to run these tests can be seen below:

```
path-to-project/LES$ bash scripts/single.sh 100
path-to-project/LES$ bash scripts/halos.sh 100
path-to-project/LES$ bash scripts/spark_cpu_standalone.sh 100 1 1
```

---

[1] http://www.dcs.gla.ac.uk/research/gpg/cluster.htm
[2] http://ark.intel.com/products/75267/Intel-Xeon-Processor-E5-2640-v2-20M-Cache-2_00-GHz
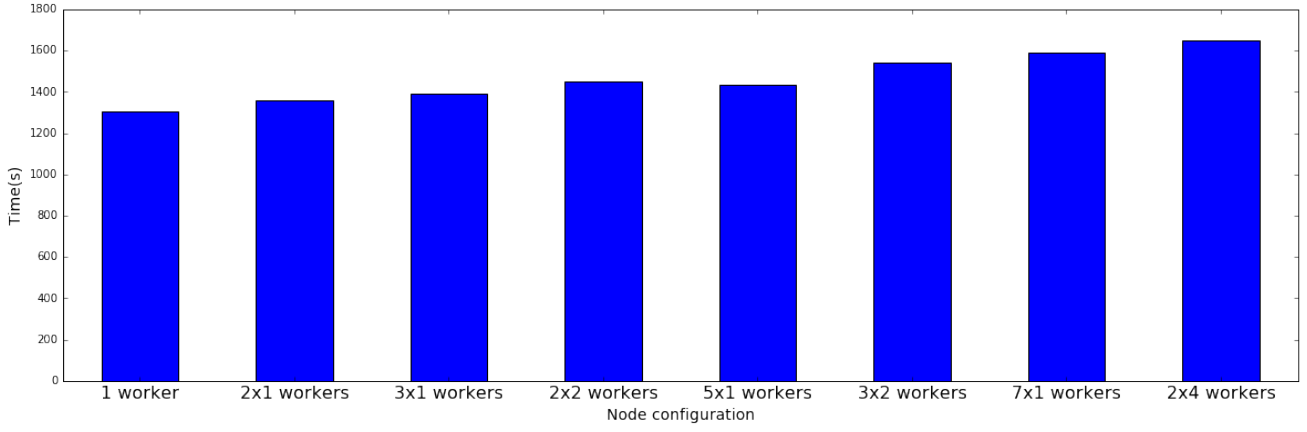
Figure 9: Bar chart with runtimes of the 100 steps of simulation for various node configurations

Test	#1 (s)	#2 (s)
Single LES	41.216	40.146
LES with simple halos	73.313	74.676
Spark local	1337.929	1354.030

Table 2: Summary of halo exchanges and reductions for each LES stage in one time step

5.3 Spark tests - increasing domain size

The final tests done using spark were done on varying number of nodes. The aim is to increase the total domain size by increasing the number of nodes. Each node runs computation on a domain of size 150x150x90, which means that the simulation on 8 nodes with configuration 2x4 will have total domain size 300x600x90. Spark master and the driver program were run on a cluster node separate from the other nodes. Each test was run only once and on 100 time steps. The results measured using the `time` command can be seen in figure 9. Commands used to run these tests can be seen below (see appendix A to find out how to run Spark master and workers):

```
path-to-project/LES$ bash scripts/spark-submit.sh 100 X Y
```

where X and Y are the sizes of the node grid.

5.4 Discussion

The tests report about 20x slowdown of the run simulation run with Spark compared to the baseline. Overhead caused by spark seems to be too high when comparing it to time spent in the simulation. In addition, a linear increase in run time with increasing number of nodes can be observed. This means that as is, the approach taken in this project might not be suitable to parallelising LES. However, due to very small number of tests, this is not a conclusive result. The next chapter goes into detail as to what tests should be done in order to make a proper conclusion.

Chapter 6

Conclusion

6.1 Future Work

While this project has proven that it is indeed possible to parallelise LES on a network cluster using Apache Spark, the first evaluation results are not favourable. However there is still some testing and evaluation to be done to determine whether using this approach to parallelisation of LES is usable.

6.1.1 Spark overhead

The biggest issue that is causing the adverse performance seems to be overhead caused by Spark. While single kernel execution, together with one halo write and one halo read takes the order of tens of milliseconds, the time it takes Spark to run the task and exchange the halos seems to be much higher. Spark reports task time to be be less than half of the overall execution time, which means that most time is spent scheduling task, resolving lineage and network transfers.

One possible explanation of the high overhead might be the need to recompute the lineage after each network shuffle, since Spark cannot predict where each item in the collection will end up after the shuffle. The fact that shuffle is necessary in each halo exchange, and that these are done more than 100 time per steps could potentially mean high frequency of Spark recalculating the lineage.

One approach to decrease the overhead could be to try different versions of Spark. Evaluation in this project was done using version 1.4.1, however since Spark seems to be rapidly moving project, there are already newer versions (1.5, 1.6), which could have better performance.

Another approach is to increase kernel execution. This could be done by increasing domain size per node. This would be decrease the ratio of Spark overhead time to execution time, which would decrease the overall slowdown from 15-20x to something more manageable. The cluster on which the evaluation was done, has 64GB of memory per node available and during the evaluation just above tenth of that has been used, so increasing the domain is a viable strategy.

6.1.2 Checkpointing

Another factor resulting in the poor performance is the frequent checkpointing (mentioned in subsection 4.4.1). The fact the the program is writing checkpoints to disk means that a lot of time is spent waiting until it is finished.

One possible solution is to write to a faster medium. In the evaluation the checkpoint directory was on a Network File System, which means that even writing to standard HDD might improve the performance. Other options are using SSD or even mounted ramfs. The best approach, however, would be to disable checkpointing entirely. While this is not possible in Spark currently, it is possible to rewrite the source code, since Spark is an open source project.

6.1.3   Evaluation on massively parallel cluster

The evaluation done in this project was done on eight nodes as the largest test. While the results show almost linear increase in computational time with increasing number of nodes, it cannot be assumed that this trend continues. Therefore it is important to test the simulation on cluster with tens, hundreds, or potentially even thousands nodes. If the computation proves to scale well on such large clusters, the current slowdown would be offset by the possibility of running the simulation on much larger domain than what is currently possible.

6.1.4   Broadcast variables

The approach proposed in section 3.2 was initially dismissed due to unnecessary network transfers. However, since it was found that the bottleneck is caused by Spark overhead and not the network transfers, this approach might prove to be surprisingly effective. This approach consists essentially only of two map steps, and thus no steps do network shuffle, which should make Spark evaluate the lineage much more efficiently. However, since this approach needs to transfer all the neighbour data to the driver program, this approach might not be usable when running the simulation on a massively parallel cluster (e.g. 1000s of nodes).

6.2   Summary

The aim of this project was to explore whether it is possible to parallelise LES using Apache Spark. The produced result proves that it is possible, but only with a large performance trade off. While it allows to run the simulation on larger domain and provides a way to nicely abstract the halo exchanges, at the current stage it does not provide suitable alternative to MPI. However, further work is needed to determine, whether the limitations discovered in this project can be overcome.

Appendices

Appendix A

Running the simulation

The project files that are required to run the parallelised version of LES can be found on GitHub. [1] The repository contains modified version of Aparapi-ucores (in `/aparapi-ucores`), the Game of Life proof of concept (in `/dummy`), and the parallelised simulation code in `/LES`.

A.1   Building Aparapi-ucores (`/aparapi-ucores`)

- Install OpenCL drivers for your system. (AMD APP SDK [2] used in project)

- Make sure environmental variable `LD_LIBRARY_PATH` contains path to OpenCL libraries

- Build Aparapi using `ant build` in `src/aparapi`. Configuration in `src/aparapi/com.amd.aparapi.jni/build.xml` might be neeeded

A.2   Running Game of life proof of concept (`/dummy`)

- Download Apache Spark[3] (project tested with version 1.4.1)

- Build project using `mvn package`

- To run the cogroup version:

```
path/to/spark/bin/spark-submit \
  --class hoos.project.dummy.DummyPartitionedJoin \
  --master local[8] \
  target/dummy-map-reduce-0.1.jar data/gol.txt data/out.txt N
```

where first two arguments are the input and ouptut files and N is the number of iterations. `local[8]` means it will run on 8 CPU threads

---

[1] https://github.com/adikus/hurricaneProject
[2] http://developer.amd.com/tools-and-sdks/opencl-zone/amd-accelerated-parallel-processing-app-sdk/
[3] http://spark.apache.org/

A.3 Running LES (`/LES`)

- Build Aparapi-ucores (above)

- Download Apache Spark

- Configure `scripts/conf.cfg`

  ```
  APARAPI_PATH=/path/to/project/aparapi-ucores/src/aparapi
  SPARK_PATH=/path/to/spark/
  IP=192.168.1.64 # IP address of Spark master
  WORK_DIR=/temp/or/scratch
  MGRID_FILE_PATH=/path/to/project/GIS/Tokyo_20mgrid.txt

  APARAPI_JNI_PATH=$APARAPI_PATH/com.amd.aparapi.jni/dist
  APARAPI_JAR_PATH=$APARAPI_PATH/com.amd.aparapi/dist/aparapi.jar
  ```

- Build project using `mvn package`

- To run version without Spark you can run either of:

  ```
  path-to-project/LES$ bash scripts/single.sh N
  path-to-project/LES$ bash scripts/halos.sh N
  ```

  where halos also runs single node halo exchanges and N is the number of iterations.

- For local mode Spark versions, run either of:

  ```
  path-to-project/LES$ bash scripts/spark_cpu_standalone.sh N X Y
  path-to-project/LES$ bash scripts/spark_gpu_standalone.sh N X Y
  ```

  to run wither on CPU or GPU and X and Y are sizes of the node grid. Note: if you specify more than 1x1 node grid, these will run sequentially.

- To run on cluster:

  - Start Spark master

    ```
    export SPARK_MASTER_IP=192.168.1.64
    sh path/to/spark/sbin/start-master.sh
    ```

    where 192.168.1.64 is replaced with the IP address configured in `conf.cfg`

  - Start one or more workers (run on the target node)

    ```
    path-to-project/LES$ bash scripts/start_cpu_worker.sh
    path-to-project/LES$ bash scripts/start_gpu_worker.sh
    ```

  - Submit spark application

    ```
    path-to-project/LES$ bash scripts/spark-submit N X Y
    ```

  - You can monitor the cluster in Spark WebUI at `http://localhost:8080/`

A.4   Creating shared library from Fortran code

The Fortran source code is not available on GitHub but in SVN repository at School of CS at University of Glasgow.

- Configure and install `https://github.com/wimvanderbauwhede/OpenCLIntegration`

- Run `scons -f SConstruct.ocl_lib` in `RefactoredSources`

- `libles_ocl.so` should have been created

Bibliography

[1] Jeffrey Dean and Sanjay Ghemawat. MapReduce: Simplified Data Processing on Large Clusters. http://static.googleusercontent.com/media/research.google.com/es/us/archive/mapreduce-osdi04.pdf, 2004. [Online; accessed 13/03/2016].

[2] Hiromasa Nakayama, Tetsuya Takemi, and Haruyasu Nagai. Les analysis of the aerodynamic surface properties for turbulent flows over building arrays with various geometries. *Journal of Applied Meteorology and Climatology*, 50(8):1692–1712, 2011.

[3] Hiromasa Nakayama, Tetsuya Takemi, and Haruyasu Nagai. Large-eddy simulation of urban boundary-layer flows by generating turbulent inflows from mesoscale meteorological simulations. *Atmospheric Science Letters*, 13(3):180–186, 2012.

[4] Siegfried Raasch and Michael Schrter. PALM A large-eddy simulation model performing on massively parallel computers. *Meteorologische Zeitschrift*, 10(5):363–372, 2001.

[5] Gordon Reid. Exploring the parallelisation of the Large Eddy Simulation using MPI and the Glasgow Model Coupling Framework, 2015.

[6] O. Segal, P. Colangelo, N. Nasiri, Zhuo Qian, and M. Margala. Aparapi-ucores: A high level programming framework for unconventional cores. In *High Performance Extreme Computing Conference (HPEC), 2015 IEEE*, pages 1–6, Sept 2015.

[7] Wim Vanderbauwhede. Model Coupling between the Weather Research and Forecasting Model and the DPRI Large Eddy Simulator for Urban Flows on GPU accelerated Multicore Systems. http://arxiv.org/abs/1504.02264, April 2015. [Online; accessed 12/03/2016].

[8] Yuzhong Yan, Lei Huang, and Liqi Yi. Is apache spark scalable to seismic data analytics and computations? In *Big Data (Big Data), 2015 IEEE International Conference on*, pages 2036–2045, Oct 2015.

[9] Matei Zaharia, Patrick Wendell, Andy Konwinski, and Holden Karau. *Learning Spark*. O'Reilly Media, Inc., February 2015. [Online; accessed 18/03/2016; https://www.safaribooksonline.com/library/view/learning-spark/9781449359034/].